

Final yaw error by encoder × adapter (rad, lower is better)

Nature CNN (scratch, 24×24)	0.189	<i>not registered</i>	0.127	<i>not registered</i>	<i>not registered</i>
Compact ViT (scratch, 14×14)	0.334	<i>not registered</i>	1.519	<i>not registered</i>	<i>not registered</i>
DINOv2 ViT-S/14 (16×16)	0.134	0.234	<i>not registered</i>	0.329	0.932
R3M ResNet-50 (layer 3, 14×14)	0.446	0.233	<i>not registered</i>	0.795	<i>not registered</i>
R3M ResNet-50 (layer 4, 7×7)	1.476	0.515	<i>not registered</i>	1.415	1.520
	Spatial softmax	Attention pooling (AFA)	Flatten	Local grid	Linear (global pool)