


What each readout draws from, across the three training approaches



One scene, stepped identically in both environments so the object pose and tabletop match; each policy is shown the input distribution it was trained on, which for VisualGoal includes the green target. Spatial-softmax rows show keypoints only (top 30% by softmax peakedness): that adapter's output is the coordinates themselves. AFA rows take the maximum over attention heads. Fields are scaled per panel and drawn over a dimmed copy of the frame