

Mean yaw error in degrees — adapter × encoder × training approach

	Nature CNN			Compact ViT			DINOv2			R3M L3			R3M L4			vs spatial softmax (shared encoders)
Attention pooling (AFA)							11	9	6	16		57	25	86	73	-23° (n=8)
Spatial softmax	11	47	8	22	83	81	8	72	23	15	83	86	87	85	89	reference
Local grid							25	35	22	53	86	86	87	86	88	+2° (n=9)
Flatten	6	89	90	88	90	82										+32° (n=6)
Linear (global pool)							47	89	82				110	86	88	+23° (n=6)
	curr	unif	vis	curr	unif	vis	curr	unif	vis	curr	unif	vis	curr	unif	vis	

curr = goal-yaw curriculum · unif = no curriculum · vis = drawn target, no curriculum. Darker is better; grey means that combination is not registered. Training metrics at iteration 6000, one run per cell.